


Chapter 4: Threads (线程)





Chapter 4: Threads

- **Overview**
- **Multithreading Models**
- **Thread Libraries**
- **Threading Issues**
- **Operating System Examples**
- **Windows XP Threads**
- **Linux Threads**





Objectives

- To introduce the notion of a thread — a fundamental unit of CPU utilization that forms the basis of multithreaded computer systems
- To discuss the APIs for the Pthreads, Win32, and Java thread libraries
- To examine issues related to multithreaded programming





4.1 Overview

- The concept of a **process** as embodying two characteristics :
 - **Unit of Resource ownership** (资源拥有单位)
 - ▶ 给每个进程分配一虚拟地址空间，保存进程映像，控制一些资源（文件，I/O设备），有状态、优先级、调度
 - **Unit of Dispatching** (调度单位)
 - ▶ 进程是由一个或多个程序的一次执行
 - ▶ 可能会与其他进程交替执行
- These two characteristics are treated **independently** by the operating system
 - Dispatching is referred to as **a thread** or **lightweight process** (轻型进程LWP)
 - Resource of ownership is referred to as **a processor task**





线程(Thread)概念

- These two characteristics are treated independently by the operating system
 - 资源拥有单元称为进程（或任务），调度的单位称为**线程**、又称**轻型进程**（light weight process）。
 - 线程只拥有一点在运行中必不可省的资源（程序计数器、一组寄存器和栈），但它可与同属一个进程的其它线程共享进程拥有的全部资源。
- **线程定义为进程内一个执行单元或一个可调度实体。**





线程

■ 线程：

- 有执行状态（状态转换）
- 不运行时保存上下文
- 有一个执行栈
- 有一些局部变量的静态存储
- 可存取所在进程的内存和其他资源
- 可以创建、撤消另一个线程





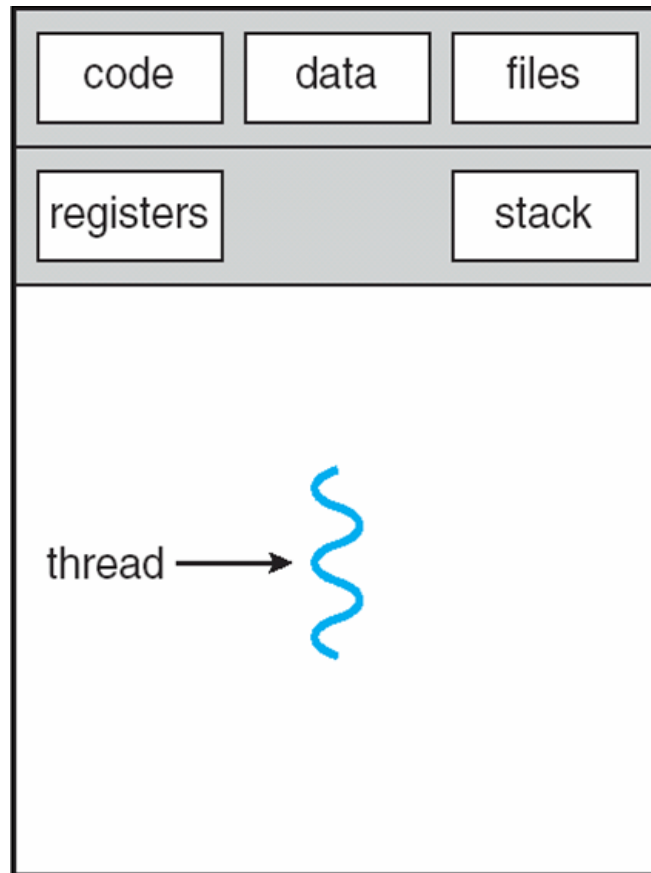
线程的特点

- 是进程的一个实体，可作为系统独立调度和分派的基本单位。
- 不拥有**系统资源**（只拥有少量的资源，资源是分配给进程）
- 一个进程中的多个线程可并发执行。（进程可创建线程执行同一程序的不同部分）
- 系统开销小、切换快。（进程的多个线程都在进程的地址空间活动）

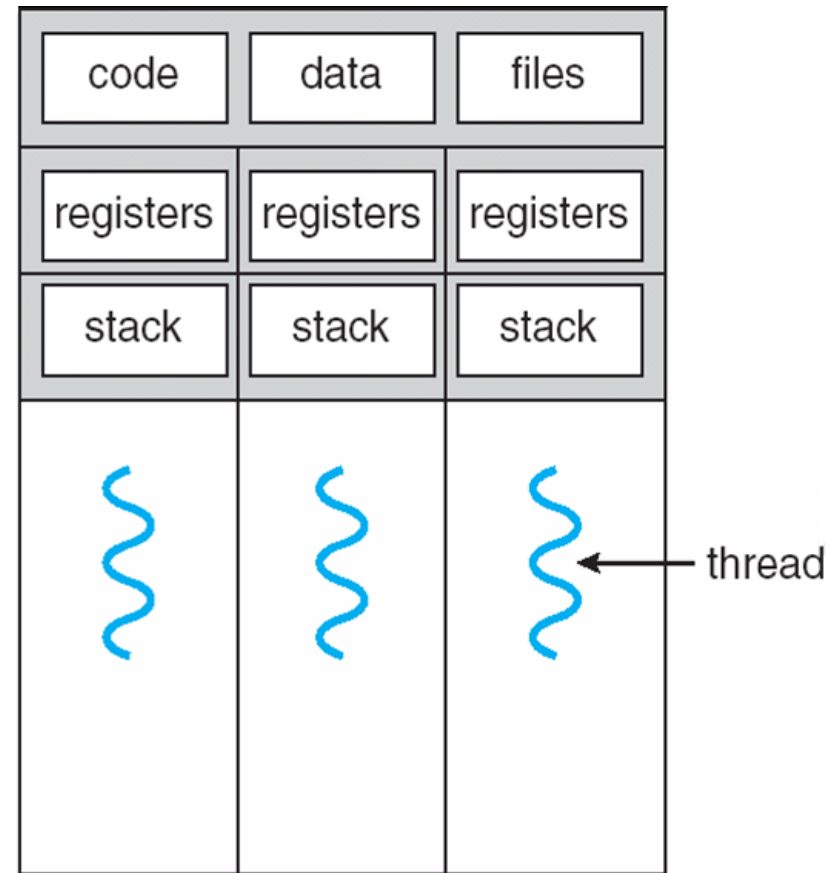




Single and Multithreaded Processes



single-threaded process



multithreaded process





Motivation (Cont.)

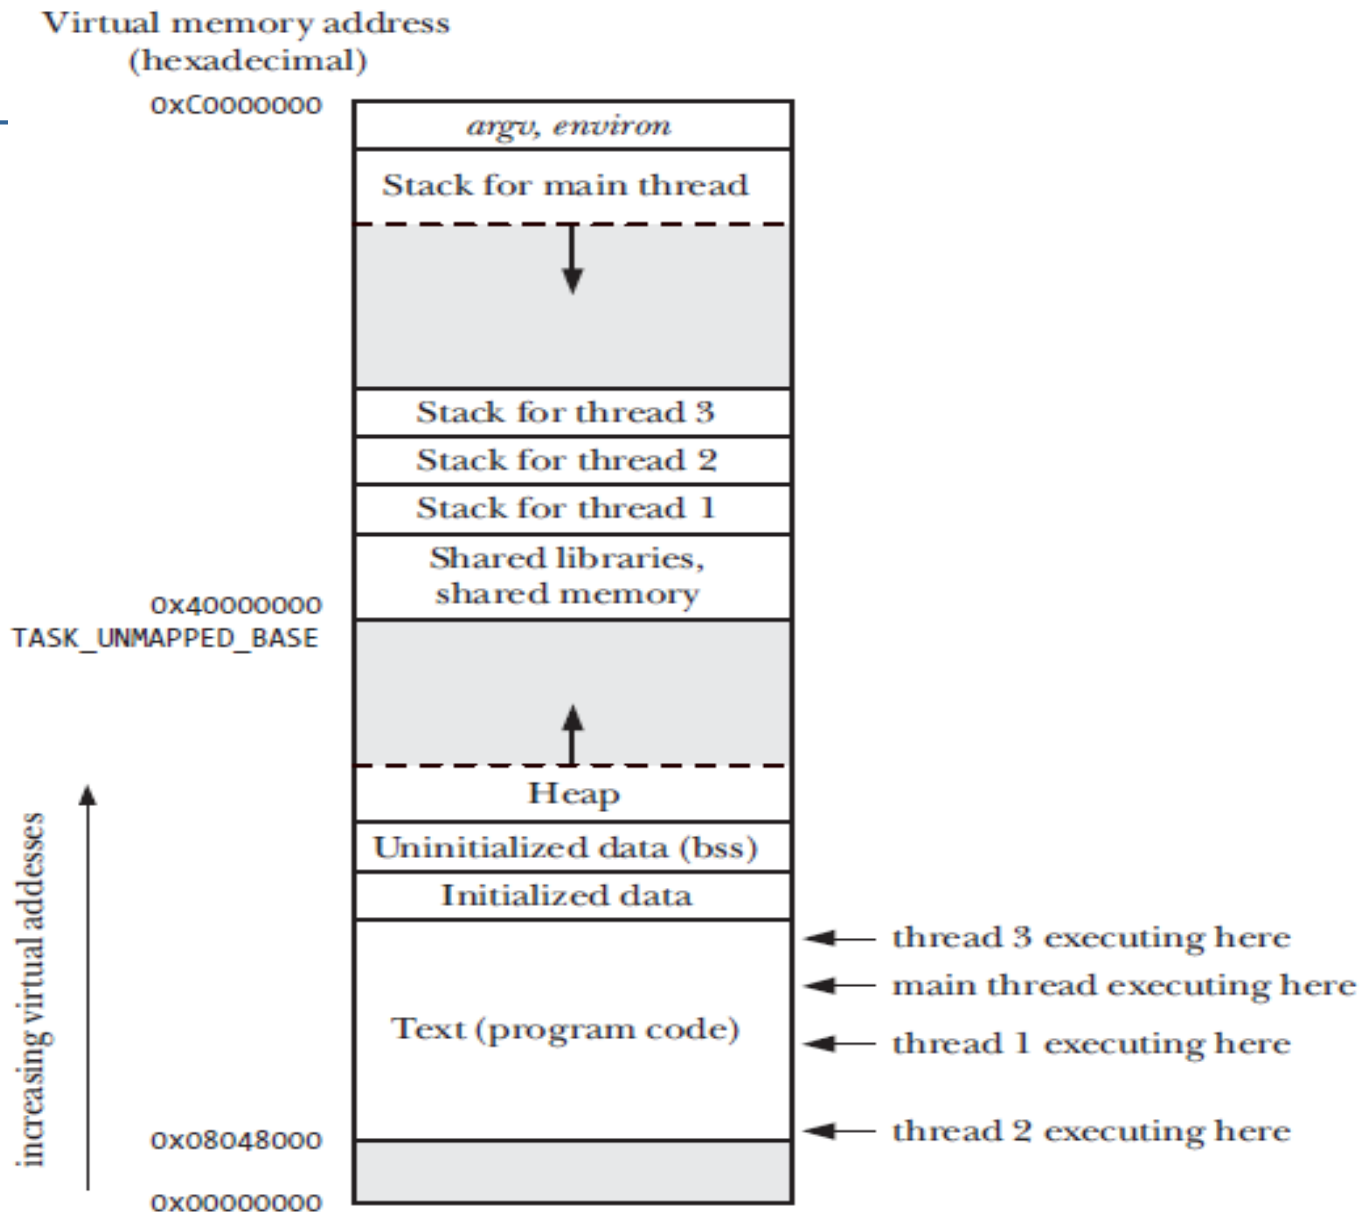
- A thread **shares with** threads belonging to the same process its:

- code section
- data section
- operating-system resources

(Process Have a virtual address space which holds the process image Protected access to processors, other processes, files, and I/O resources)

- A **traditional** or **heavyweight process** (重型线程) is equal to a task with one thread





Four threads executing in a process (Linux/x86-32)





Benefits

- 创建一个新线程花费时间少（结束亦如此）
- 两个线程的切换花费时间少
(如果机器设有“存储[恢复]所有寄存器”指令，则整个切换过程用几条指令即可完成)
- 因为同一进程内的线程共享内存和文件，因此它们之间相互通信无须调用内核
- 适合多处理机系统





例子1:

- LAN中的一个文件服务器，在一段时间内需要处理几个文件请求
 - 有效的方法是：为每一个请求创建一个线程
 - 在一个SMP机器上：多个线程可以同时在不同的处理器上运行





例子2:

- 一个线程显示菜单，并读入用户输入；另一个线程执行用户命令
 - 考虑一个应用：由几个独立部分组成，这几个部分不需要顺序执行，则每个部分可以以线程方式实现
 - 当一个线程因I/O阻塞时，可以切换到同一应用的另一个线程

Example: [createthread.cpp](#)

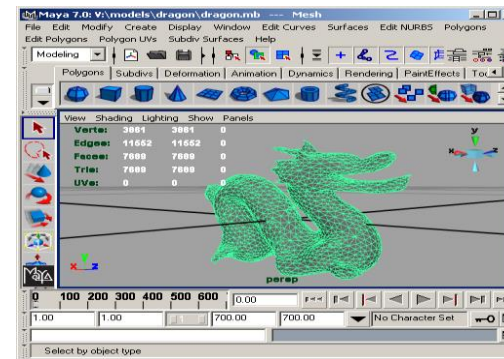




Multicore Programming

■ **Multicore systems** putting pressure on programmers, challenges include

- Dividing activities
- Balance
- Data splitting
- Data dependency
- Testing and debugging




Each
can run
on its
own
core






Multithreaded Server Architecture






Concurrent Execution on a Single-core System





Parallel Execution on a Multicore System





线程的实现机制

- 用户级线程 user-level thread
- 核心级线程 kernel-level thread
- 两者结合方法





User Threads（用户级线程）

- **用户级线程**：不依赖于OS核心（内核不了解用户线程的存在），应用进程利用线程库提供创建、同步、调度和管理线程的函数来控制用户线程。

如：数据库系统informix，图形处理Aldus PageMaker。调度由应用软件内部进行，通常采用非抢先式和更简单的规则，也无需用户态/核心态切换，所以速度特别快。一个线程发起系统调用而阻塞，则整个进程在等待。

- 用户线程的维护由应用进程完成；
- 内核不了解用户线程的存在；
- 用户线程切换不需要内核特权；
- 用户线程调度算法可针对应用优化；
- 一个线程发起系统调用而阻塞，则整个进程在等待。

- Three primary **thread libraries**:

- POSIX **Pthreads** 、 Win32 threads、 Java threads





Kernel Threads (内核级线程)

■ **内核级线程**：依赖于OS核心，由内核的内部需求进行创建和撤销，用来执行一个指定的函数。一个线程发起系统调用而阻塞，不会影响其他线程。时间片分配给线程，所以多线程的进程获得更多CPU时间。

- 内核维护进程和线程的上下文信息；
- 线程切换由内核完成；
- 时间片分配给线程，所以多线程的进程获得更多CPU时间；
- 一个线程发起系统调用而阻塞，不会影响其他线程的运行。

■ Examples

- Windows XP/2000 及以后
- Solaris
- Linux
- POSIX Pthreads
- Mac OS X





4.2 Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many






Many-to-One

- Many user-level threads mapped to single kernel thread
- Implemented by user-level runtime libraries
 - Create, schedule, synchronize threads at user-level
- OS is not aware of user-level threads
 - OS thinks each process contains only a single thread of control
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads





Many-to-One Model



■ Advantages

- Does not require OS support
- Can tune scheduling policy to meet application (user level) demands
- Lower overhead thread operations since no system calls

■ Disadvantages

- Cannot leverage multiprocessors (no true parallelism)
- **Entire process blocks when one thread blocks**






One-to-One

- **Each user-level thread maps to kernel thread**
- **OS provides each user-level thread with a kernel thread**
- **Each kernel thread scheduled independently**
- **Thread operations (creation, scheduling, synchronization) performed by OS**
- **Examples**
 - **Windows NT/XP/2000**
 - **Linux**
 - **Solaris 9 and later**





One-to-one Model



■ Advantages

- Each kernel-level thread can run in **parallel on a multiprocessor**
- When one thread blocks, **other threads from process can be scheduled**

■ Disadvantages

- **Higher overhead** for thread operations
- OS must scale well with increasing number of threads






Many-to-Many Model

- **Allows many user level threads to be mapped to many kernel threads**
- **Allows the operating system to create a sufficient number of kernel threads**
- **Examples:**
 - Solaris prior to version 9**
 - Windows NT/2000 with the ThreadFiber**





Many-to-Many Model






Two-level Model

- **Similar to M:M, except that it allows a user thread to be bound to kernel thread**
- **Examples**
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier








Two-level Model





Solaris多线程模型的发展 - M:N -> 1:1



LWP

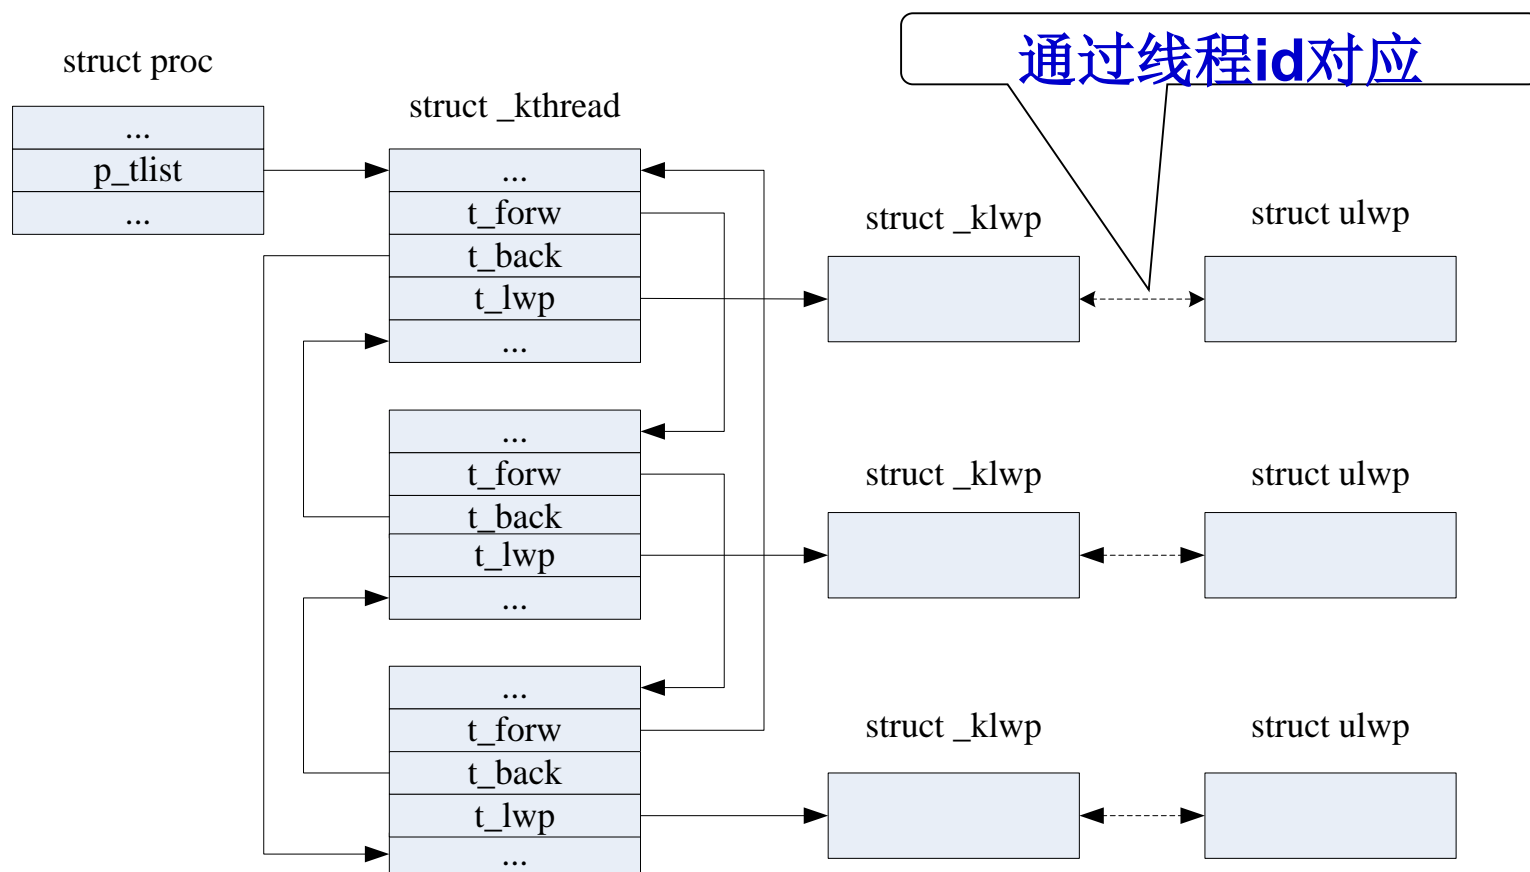


Kernel thread





Solaris用户线程、内核线程、lwp三者之间的关系





4.3 Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in **user space**
 - **Kernel-level library** supported by the OS





Pthreads

- May be provided either as **user-level or kernel-level**
- A **POSIX** (Portable Operating System Interface) standard (IEEE 1003.1c) API for thread creation and synchronization

http://standards.ieee.org/reading/ieee/stad_public/description/posix

- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (**Solaris, Linux, Mac OS X**)





Java Threads

- **Java threads are managed by the JVM**
- **Typically implemented using the threads model provided by underlying OS**
- **Java threads may be created by:**
 - **Extending Thread class**
 - **Implementing the Runnable interface**





4.4 Threading Issues

- Semantics of fork() and exec() system calls
- Thread cancellation of target thread
 - Asynchronous or deferred
- Signal handling
- Thread pools
 - Thread-specific data
- Scheduler activations





Semantics of `fork()` and `exec()`

`fork()`

- duplicate only the calling thread or all threads?

`exec()`

- Replaces the process - including all threads?





Thread Cancellation

- Terminating a thread before it has finished
- Two general approaches:
 - **Asynchronous** cancellation terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled





Signal Handling (信号处理)

- Signals are used in UNIX systems to notify a process that a particular event has occurred
- A **signal handler** is used to process signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Signal is handled
- Options:
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process





Thread Pools (线程池)

- **Create a number of threads in a pool where they await work**
- **Advantages:**
 - **Usually slightly faster to service a request with an existing thread than create a new thread**
 - **Allows the number of threads in the application(s) to be bound to the size of the pool**





Thread Specific Data (线程特有数据)

- Allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)





Scheduler Activations

- Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the thread library
- This communication allows an application to maintain the correct number kernel threads





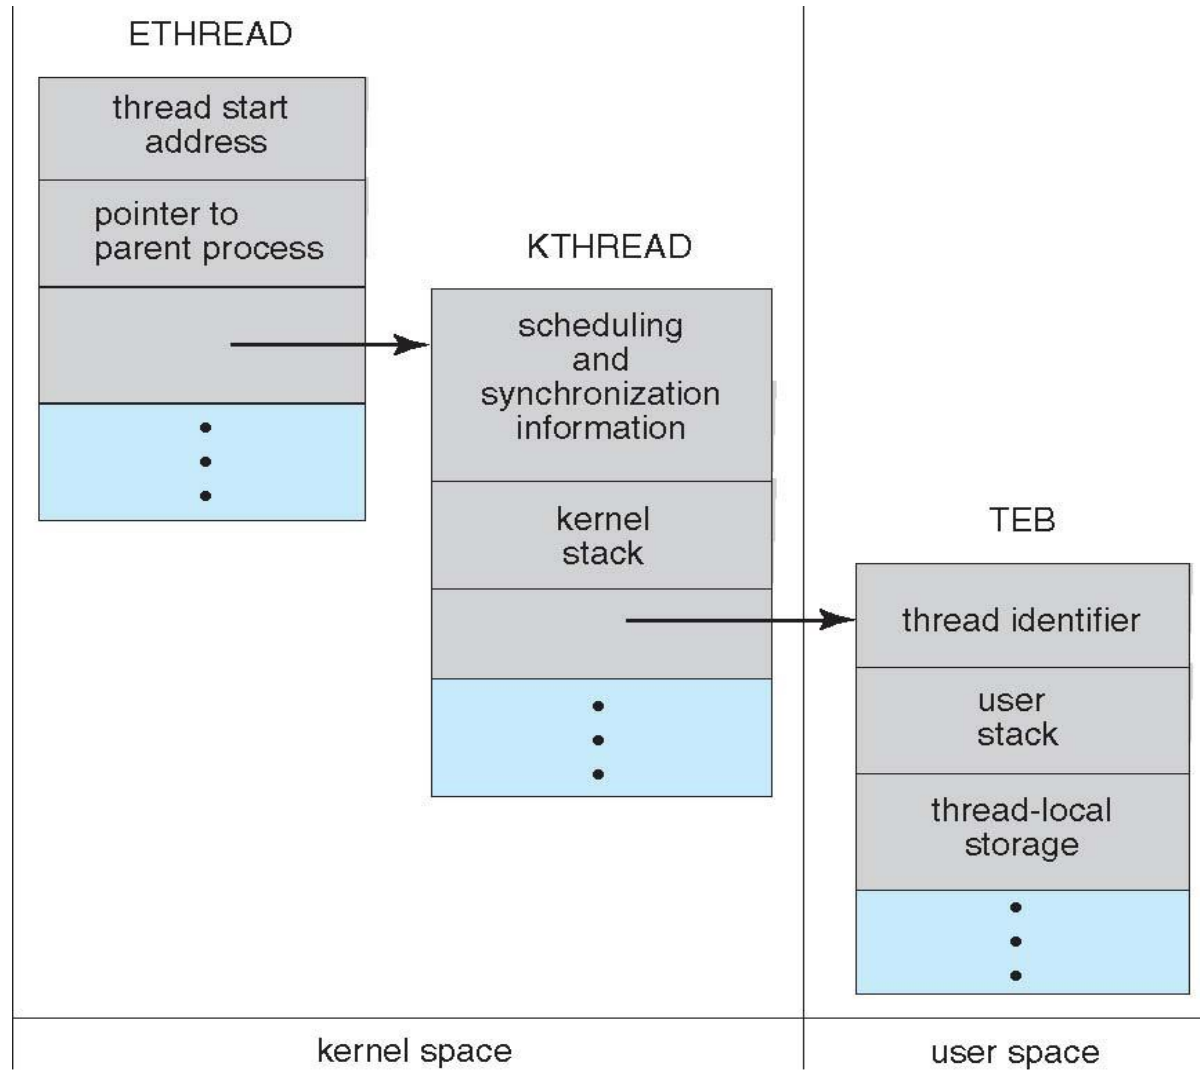
Operating System Examples

- **Windows XP Threads**
- **Linux Thread**





Windows XP Threads





Windows XP Threads

- Implements the **one-to-one** mapping, kernel-level
- Each thread contains
 - A thread id
 - Register set
 - Separate user and kernel stacks
 - Private data storage area
- The register set, stacks, and private storage area are known as the **context** of the threads
- The primary data structures of a thread include:
 - ETHREAD (executive thread block)
 - KTHREAD (kernel thread block)
 - TEB (thread environment block)

CreateThread.cpp





Linux Threads

- Linux refers to them as **tasks** rather than *threads*
- Thread creation is done through **clone()** system call
- **clone()** allows a child task to share the address space of the parent task (process)





Linux Threads

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.





Exercise

■ ex_ch3-4.doc





Homework

- 按时完成“作业系统”的作业
4.4、4.7






Reading Assignments

■ Read for this week:

- **Chapters 4**
of the text book:

■ Read for next week:

- **Chapters 5**
of the text book:



End of Chapter 4

